

Design and Analysis of Algorithms

FALL-25

Project



Sagar Mehmood (23i-0562)

Warisha Shaukat (23i-0809)

Sarah Suhail (23i-0557)

CS-5B

Performance Analysis of Binary, Fibonacci, and Hollow Heaps in Shortest Path Problems

1. Theoretical Overview

1.1 Binary Heap

- The Binary Heap serves as the standard baseline for priority queue implementations. Structurally, it is a complete binary tree where every node satisfies the heap invariant: a parent's key is always less than or equal to that of its children. Rather than using explicit pointers, the tree is typically flattened into a dynamic array, where children of a node at index i , i are located at $2i + 1$ and $2i + 2$. This implicit mapping eliminates the need to store structural pointers.
- **Advantages:** The array-based layout offers exceptional memory efficiency and cache locality, as traversing the tree involves predictable, contiguous memory access patterns.
- **Disadvantages:** Operations that require rearranging the tree structure, such as `decrease_key` and `merge`, are relatively expensive ($O(\log n)$ and $O(n)$, respectively) compared to more sophisticated pointer-based heaps.

1.2 Fibonacci Heap

- Introduced by Fredman and Tarjan, the Fibonacci Heap is a sophisticated pointer-based structure designed to theoretically accelerate graph algorithms. Instead of a single rigid tree, it maintains a forest of heap-ordered trees. Its defining characteristic is "laziness": operations like `insert` and `decrease_key` simply add nodes or cut subtrees without immediately restructuring the heap. The heavy lifting—consolidating trees by rank—is deferred until an `extract_min` operation is performed.
- **Advantages:** It achieves a theoretically optimal amortized time complexity of $O(1)$ for `decrease_key`, making it the gold standard for algorithms like Dijkstra's and Prim's on dense graphs.
- **Disadvantages:** The structure requires significant memory overhead (each node needs pointers for parent, child, and siblings, plus degree and mark bits). In practice, the high constant factors and complex pointer manipulation often make it slower than simpler heaps for all but the largest inputs.

1.3 Hollow Heap

- The Hollow Heap, proposed by Hansen and Tarjan in 2015, offers a modern alternative that simplifies the implementation of asymptotically optimal heaps. It achieves bounds similar to the Fibonacci heap ($O(1)$ amortized `decrease_key`) but abandons the complex "cascading cut" logic. Instead, it employs a lazy deletion strategy: when a key is decreased, the old

node is not moved but simply marked as "hollow" (effectively deleted) and left in place, while a new node with the updated key is inserted. These hollow nodes are only removed during the `extract_min` cleanup phase.

- **Advantages:** It is significantly easier to implement than the Fibonacci heap while retaining the same theoretical efficiency. The logic for linking and ranking trees is streamlined around the concept of hollow nodes.
- **Disadvantages:** The lazy deletion approach can lead to a larger memory footprint, as "zombie" hollow nodes persist in the structure until the minimum extraction process sweeps them away.

1.4 Asymptotic Complexity Comparison

Operation	Binary Heap (Worst Case)	Fibonacci Heap (Amortized)	Hollow Heap (Amortized)
Insert	$O(\log N)$	$O(1)$	$O(1)$
Extract Min	$O(\log N)$	$O(\log N)$	$O(\log N)$
Decrease Key	$O(\log N)$	$O(1)$	$O(1)$

2. Implementation Details

- The project was implemented in C++ using an Object-Oriented approach.
- Interface: A template abstract base class `PriorityQueue` defined the contract, ensuring Dijkstra's algorithm remained agnostic to the underlying data structure.
- Graph Storage: An Adjacency List (`:std::vector<std::vector<GraphEdge>>`) was used to store the road networks efficiently ($O(V+E)$ space).
- Optimization: The Binary Heap utilized an intrusive index within the node handle to allow $O(1)$ lookup of nodes during `decrease_key`, preventing the standard $O(N)$ linear scan.

3. Experimental Setup

- Algorithm: Dijkstra's Single Source Shortest Path
- Datasets: Real-world road networks of varying sizes.
- Hongkong: ~43k Nodes, ~91k Edges (Small)
- Shanghai: ~390k Nodes, ~855k Edges (Medium)
- Chongqing: ~1.2M Nodes, ~2.4M Edges (Large)

Metrics Recorded:

- Operation Latency: Average time for Insert, Extract-Min, and Decrease-Key operations in microseconds.

- Total Runtime: Total execution time for Dijkstra's algorithm in milliseconds (ms).
- Structural Metrics: Maximum Tree Height, Maximum Number of Roots (for Fibonacci/Hollow heaps), Total Link Operations, Consolidation Passes, and Peak Memory Usage (MB).

4. Results & Analysis

4.1 Performance Data

Dataset 1: Hongkong (Small)

== Batch Summary for Hong Kong road network ==									
Heap	Runtime(ms)	Inserts	Insert Avg (us)	Extracts	Extract Avg (us)	Decreases	Decrease Avg (us)	Reachable	
Binary	15	43620	0.061	43620	0.088	605	0.034	43620	
Fibonacci	29	43620	0.047	43620	0.455	601	0.025	43620	
Hollow	19	43620	0.107	43620	0.123	602	0.116	43620	
== Structural Metrics for Hong Kong road network ==									
Heap	MaxNodes	MaxBytes	MaxMB	Height	MaxRoots	ConsolPasses	LinkOps		
Binary	100	2728	0.003	7	1	43564	213704		
Fibonacci	98	6272	0.006	8	14	43612	185072		
Hollow	99	2825392	2.695	7	1	43620	329612		


Dataset 2: Shanghai (Medium)

== Batch Summary for Shanghai road network ==									
Heap	Runtime(ms)	Inserts	Insert Avg (us)	Extracts	Extract Avg (us)	Decreases	Decrease Avg (us)	Reachable	
Binary	141	390171	0.045	390171	0.084	14261	0.028	390171	
Fibonacci	329	390171	0.045	390171	0.575	14256	0.026	390171	
Hollow	207	390171	0.108	390171	0.177	14248	0.074	390171	
== Structural Metrics for Shanghai road network ==									
Heap	MaxNodes	MaxBytes	MaxMB	Height	MaxRoots	ConsolPasses	LinkOps		
Binary	662	16280	0.016	10	1	390036	2825257		
Fibonacci	662	42368	0.040	11	20	390166	2579158		
Hollow	663	25768720	24.575	10	1	390171	4403098		


Dataset 3: Chongqing (Large)

== Batch Summary for Chongqing road network ==									
Heap	Runtime(ms)	Inserts	Insert Avg (us)	Extracts	Extract Avg (us)	Decreases	Decrease Avg (us)	Reachable	
Binary	425	1185464	0.043	1185464	0.075	9921	0.026	1185464	
Fibonacci	922	1185464	0.043	1185464	0.517	9927	0.027	1185464	
Hollow	562	1185464	0.108	1185464	0.135	9935	0.079	1185464	
== Structural Metrics for Chongqing road network ==									
Heap	MaxNodes	MaxBytes	MaxMB	Height	MaxRoots	ConsolPasses	LinkOps		
Binary	499	13672	0.013	9	1	1185372	7308290		
Fibonacci	499	31936	0.030	11	19	1185449	6689216		
Hollow	498	76426056	72.886	9	1	1185464	11657073		


RunTime Analysis :



Decrease Operation Analysis :



Memory Usage :



4.2 Analysis Discussion

1. Operation Speed:

- **Decrease-Key:** Contrary to theoretical expectations, the Binary Heap (0.026 μ s) matched or outperformed the Fibonacci Heap (0.027 μ s) and the Hollow Heap (0.079 μ s). While advanced heaps provide O(1) amortized decrease-key, the constant factors—especially memory allocation and pointer chasing—proved far more significant in practice. The Hollow Heap was notably slower because every decrease-key operation requires allocating a new node.
- **Extract-Min:** The cost of “lazy” heap structures became evident here. The Fibonacci Heap was roughly seven times slower (0.517 μ s) than the Binary Heap (0.075 μ s) due to the heavy consolidation required during extract-min. The Hollow Heap performed better (0.135 μ s) but was still almost twice as slow as the Binary Heap.

2. Total Runtime:

The Binary Heap achieved the fastest overall runtime (425 ms), followed by the Hollow Heap (562 ms), while the Fibonacci Heap was significantly slower (922 ms).

Reasoning: Road networks are sparse graphs ($E \approx 2V$), resulting in a relatively low number of decrease-key operations (around 9,900) compared to insert and extract

operations (around 1.2 million). The Binary Heap's contiguous memory layout provided superior cache locality, whereas the pointer-heavy structures of the Fibonacci and Hollow heaps resulted in frequent cache misses.

3. Memory Overhead:

- **Fibonacci Heap:** Consumed approximately 2.3× more memory (0.030 MB) than the Binary Heap (0.013 MB), mainly due to storing multiple pointers (parent, child, left sibling, right sibling) per node.
 - **Hollow Heap:** Consumed dramatically more memory (72.89 MB)—over 5,000× that of the Binary Heap. This results directly from its “hollow node” design: old nodes are not removed immediately but remain in the DAG. During long-running processes like Dijkstra on a large graph, these unused “zombie” nodes accumulate, creating severe memory overhead.
-

5. Conclusion

Although Fibonacci and Hollow heaps provide better theoretical asymptotic performance for dense graphs with frequent key updates, this study clearly shows that the Binary Heap is practically superior for real-world road-network graphs.

The theoretical advantage of $O(1)$ decrease-key was neutralized by:

1. **Sparsity:** Real-world road networks have a very low number of decrease-key operations relative to other operations.
2. **System-Level Costs:** Dynamic memory allocation (Hollow Heap) and pointer manipulation (Fibonacci Heap) introduce significant overhead compared to the cache-friendly array layout of the Binary Heap.

Recommendation: For the logistics company's routing system, the Binary Heap is the preferred choice. It achieved the fastest runtime, the lowest memory usage by a huge margin, and has the simplest implementation.